

The Essence of C++

with examples in C++84, C++98, C++11, and C++14

Bjarne Stroustrup


Texas A&M University

www.stroustrup.com




Overview

- Aims and constraints
- C++ in four slides
- Resource management
- OOP: Classes and Hierarchies
 - (very briefly)
- GP: Templates
 - Requirements checking
- Challenges




What did/do I want?

- Type safety
 - Encapsulate necessary unsafe operations
- Resource safety
 - It's not all memory
- Performance
 - For some parts of almost all systems, it's important
- Predictability
 - For hard and soft real time
- Teachability
 - Complexity of code should be proportional to the complexity of the task
- Readability
 - People and machines (“analyzability”)



Who did/do I want it for?

- Primary concerns
 - Systems programming
 - Embedded systems
 - Resource constrained systems
 - Large systems
- Experts
 - “C++ is expert friendly”
- Novices
 - C++ is not *just* expert friendly



What is C++?

Template
meta-programming!

Class hierarchies

A hybrid language

A multi-paradigm
programming language

It's C!

Embedded systems
programming language

Low level!

A random collection
of features

Generic programming

An object-oriented
programming language

Too big!

Classes

Buffer
overflows



C++

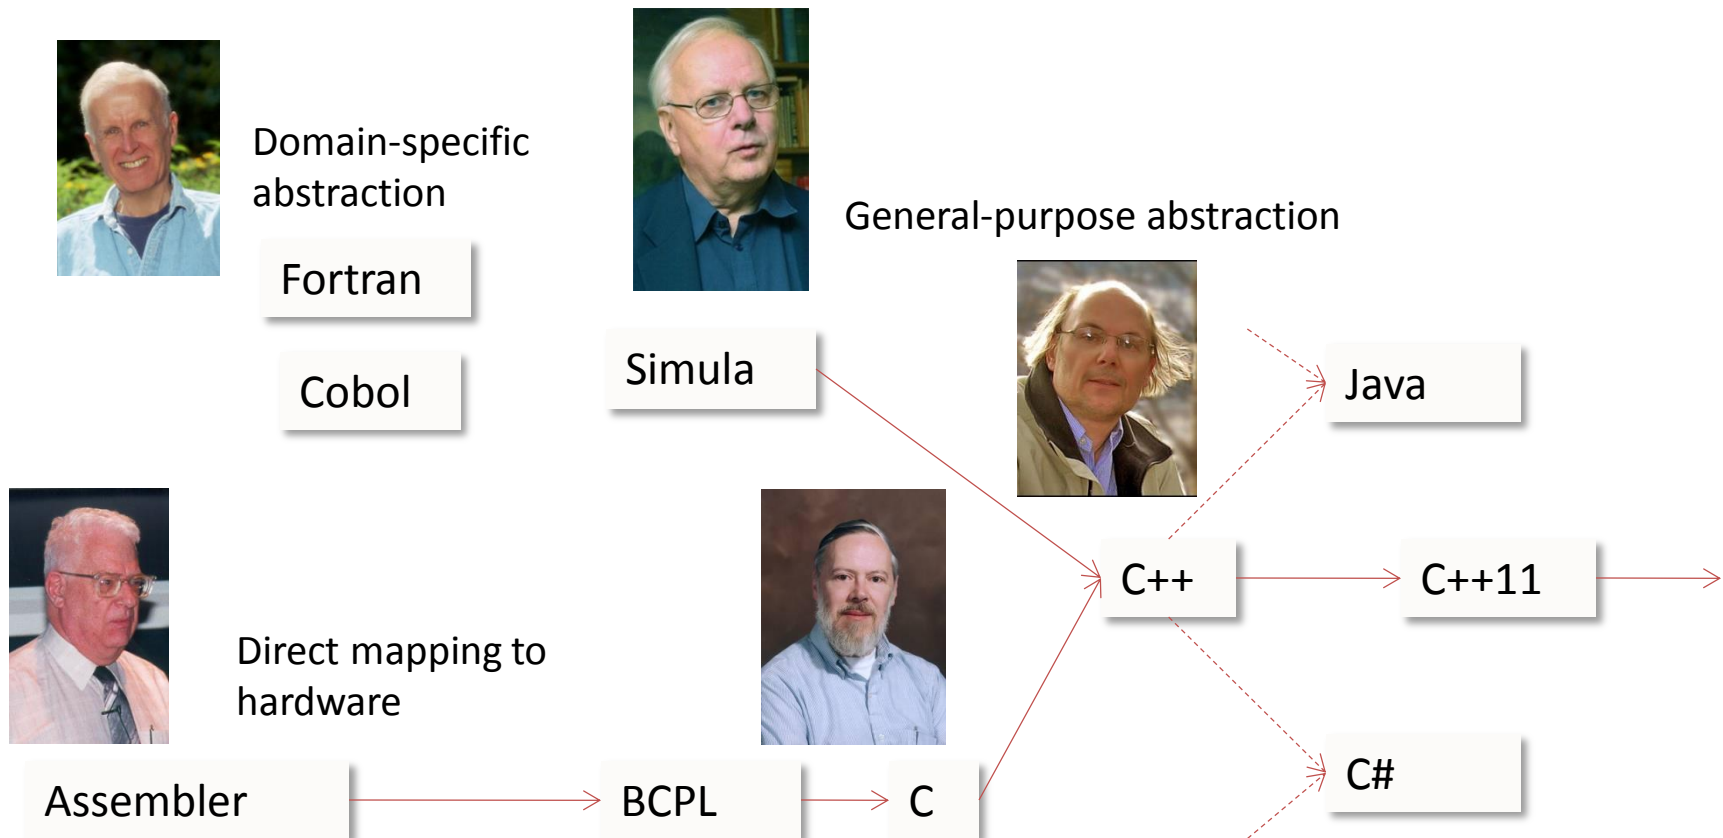
A light-weight abstraction programming language



Key strengths:

- software infrastructure
- resource-constrained applications

Programming Languages




What does C++ offer?

- Not perfection
 - Of course
- Not everything for everybody
 - Of course
- A solid fundamental model
 - Yes, really
- 30+ years of real-world “refinement”
 - It works
- Performance
 - A match for anything
- The best is buried in “compatibility stuff”
 - long-term stability is a feature



What does C++ offer?

- C++ in Four slides
 - Map to hardware
 - Classes
 - Inheritance
 - Parameterized types



- If you understand **int** and **vector**, you understand C++
 - The rest is “details” (1,300+ pages of details)

Map to Hardware

- Primitive operations => instructions

- +, %, ->, [], (), ...

value

- **int**, double, complex<double>, Date, ...

handle



- **vector**, string, thread, Matrix, ...

value

- Objects can be composed by simple concatenation:

- Arrays

- Classes/structs



Classes: Construction/Destruction

- From the first week of “C with Classes” (1979)

```
class X {                // user-defined type
public:                  // interface
    X(Something);        // constructor from Something
    ~X();                // destructor
    // ...
private:                // implementation
    // ...
};
```



“A constructor establishes the environment for the members to run in; the destructor reverses its actions.”

Abstract Classes and Inheritance

- Insulate the user from the implementation

```
struct Device {                                // abstract class  
    virtual int put(const char*) = 0;         // pure virtual function  
    virtual int get(const char*) = 0;  
};
```

- No data members, all data in derived classes
 - “not brittle”
- Manipulate through pointer or reference
 - Typically allocated on the free store (“dynamic memory”)
 - Typically requires some form of lifetime management (use resource handles)
- Is the root of a hierarchy of derived classes

Parameterized Types and Classes

- Templates
 - Essential: Support for generic programming
 - Secondary: Support for compile-time computation

```
template<typename T>
```

```
class vector { /* ... */ };           // a generic type
```

```
vector<double> constants = {3.14159265359, 2.54, 1, 6.62606957E-34, }; // a use
```

```
template<typename C>
```

```
void sort (Cont& c) { /* ... */ }     // a generic function
```

```
sort(constants);                     // a use
```



Not C++ (fundamental)

- No crucial dependence on a garbage collector
 - GC is a last and imperfect resort
- No guaranteed type safety
 - Not for all constructs
 - C compatibility, history, pointers/arrays, unions, casts, ...
- No virtual machine
 - For many reasons, we often want to run on the real machine
 - You can run on a virtual machine (or in a sandbox) if you want to



Not C++ (market realities)

- No huge “standard” library
 - No owner
 - To produce “free” libraries to ensure market share
 - No central authority
 - To approve, reject, and help integration of libraries
- No standard
 - Graphics/GUI
 - Competing frameworks
 - XML support
 - Web support
 - ...



Resource Management




Resource management

- A resource should be owned by a “handle”
 - A “handle” should present a well-defined and useful abstraction
 - E.g. a vector, string, file, thread
- Use constructors and a destructor

```
class Vector {                                // vector of doubles
    Vector(initializer_list<double>); // acquire memory; initialize elements
    ~Vector();                          // destroy elements; release memory
    // ...
private:
    double* elem;    // pointer to elements
    int sz;          // number of elements
};

void fct()
{
    Vector v {1, 1.618, 3.14, 2.99e8}; // vector of doubles
    // ...
}
```



Resource management

- A handle usually is scoped
 - Handles lifetime (initialization, cleanup), and more

```
Vector::Vector(initializer_list<double> lst)
    :elem {new double[lst.size()]}, sz{lst.size()};    // acquire memory
{
    uninitialized_copy(lst.begin(),lst.end(),elem);    // initialize elements
}
```

```
Vector::~~Vector()
{
    delete[] elem;    // destroy elements; release memory
};
```


Resource management

- What about errors?
 - A resource is something you acquire and release
 - A resource should have an owner
 - Ultimately “root” a resource in a (scoped) handle
 - “Resource Acquisition Is Initialization” (RAII)
 - Acquire during construction
 - Release in destructor
 - Throw exception in case of failure
 - Can be simulated, but not conveniently
 - Never throw while holding a resource **not** owned by a handle
- In general
 - Leave established invariants intact when leaving a scope

“Resource Acquisition is Initialization” (RAII)

- For all resources
 - Memory (done by `std::string`, `std::vector`, `std::map`, ...)
 - Locks (e.g. `std::unique_lock`), files (e.g. `std::fstream`), sockets, threads (e.g. `std::thread`), ...

```
std::mutex mtx;      // a resource  
int sh;              // shared data
```

```
void f()  
{  
    std::lock_guard lck {mtx}; // grab (acquire) the mutex  
    sh+=1;                    // manipulate shared data  
}                             // implicitly release the mutex
```

Pointer Misuse

- Many (most?) uses of pointers in local scope are not exception safe

```
void f(int n, int x)
{
    Gadget* p = new Gadget{n};           // look I'm a java programmer! 😊
    // ...
    if (x<100) throw std::runtime_error{"Weird!"};           // leak
    if (x<200) return;                                       // leak
    // ...
    delete p;        // and I want my garbage collector! ☹️
}
```

- But, garbage collection would not release non-memory resources anyway
- But, why use a “naked” pointer?

Resource Handles and Pointers

- A `std::shared_ptr` releases its object at when the last `shared_ptr` to it is destroyed

```
void f(int n, int x)
{
    shared_ptr<Gadget> p {new Gadget{n}};    // manage that pointer!
    // ...
    if (x<100) throw std::runtime_error{"Weird!"};    // no leak
    if (x<200) return;                                // no leak
    // ...
}
```

- `shared_ptr` provides a form of garbage collection
- But I'm not sharing anything
 - use a `unique_ptr`

Resource Handles and Pointers

- But why use a pointer at all?
- If you can, just use a scoped variable

```
void f(int n, int x)
{
    Gadget g {n};
    // ...
    if (x<100) throw std::runtime_error{"Weird!"};    // no leak
    if (x<200) return;                                // no leak
    // ...
}
```


Why do we use pointers?

- And references, iterators, etc.
- To represent ownership
 - **Don't!** Instead, use handles
- To reference resources
 - from within a handle
- To represent positions
 - Be careful
- To pass large amounts of data (into a function)
 - E.g. pass by **const** reference
- To return large amount of data (out of a function)
 - **Don't!** Instead use move operations

How to get a lot of data cheaply out of a function?

- Ideas
 - Return a pointer to a **new**'d object
 - Who does the **delete**?
 - Return a reference to a **new**'d object
 - Who does the **delete**?
 - Delete what?
 - Pass a target object
 - We are regressing towards assembly code
 - Return an object
 - Copies are expensive
 - Tricks to avoid copying are brittle
 - Tricks to avoid copying are not general
 - Return a handle
 - Simple and cheap

Move semantics

- Return a **Matrix**

```
Matrix operator+(const Matrix& a, const Matrix& b)
```

```
{
```

```
    Matrix r;
```


```
    // copy a[i]+b[i] into r[i] for each i
```

```
    return r;
```

```
}
```

```
Matrix res = a+b;
```

- Define move a constructor for **Matrix**
 - don't copy; “steal the representation”




Move semantics

- Direct support in C++11: Move constructor

```
class Matrix {  
    Representation rep;  
    // ...  
    Matrix(Matrix&& a)           // move constructor  
    {  
        rep = a.rep;           // *this gets a's elements  
        a.rep = {};            // a becomes the empty Matrix  
    }  
};
```

Matrix res = a+b;



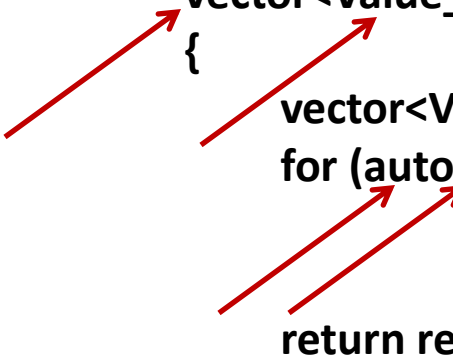
No garbage collection needed

- For general, simple, implicit, and efficient resource management
- Apply these techniques in order:
 1. Store data in containers
 - The semantics of the fundamental abstraction is reflected in the interface
 - Including lifetime
 2. Manage **all** resources with resource handles
 - RAI
 - Not just memory: **all** resources
 3. Use “smart pointers”
 - They are still pointers
 4. Plug in a garbage collector
 - For “litter collection”
 - C++11 specifies an interface
 - Can still leak non-memory resources

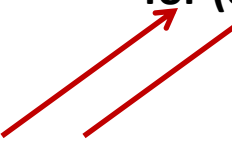
Range-for, auto, and move

- As ever, what matters is how features work in combination

```
template<typename C, typename V>
vector<Value_type<C>*> find_all(C& c, V v) // find all occurrences of v in c
{
    vector<Value_type<C>*> res;
    for (auto& x : c)
        if (x==v)
            res.push_back(&x);
    return res;
}
```



```
string m {"Mary had a little lamb"};
for (const auto p : find_all(m,'a')) // p is a char*
    if (*p!='a')
        cerr << "string bug!\n";
```



RAII and Move Semantics

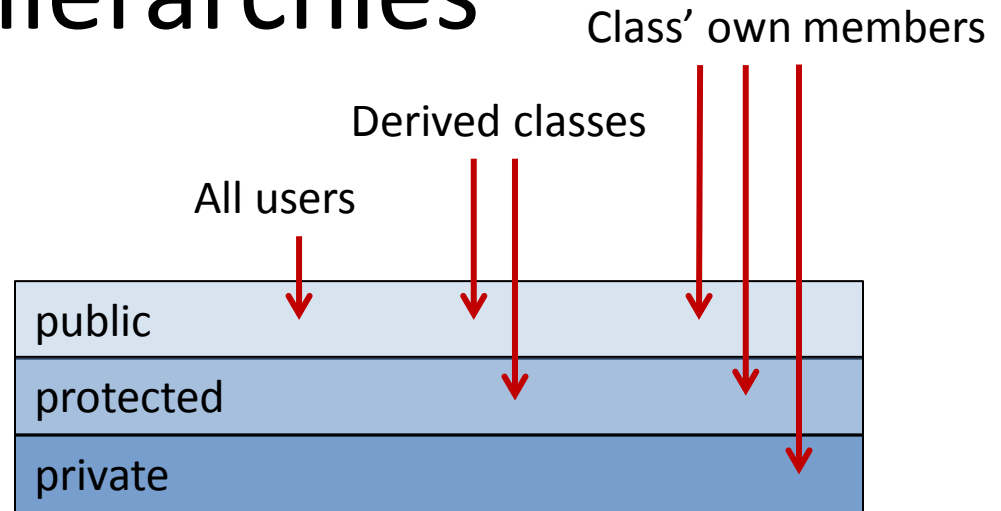
- All the standard-library containers provide it
 - **vector**
 - **list, forward_list** (singly-linked list), ...
 - **map, unordered_map** (hash table),...
 - **set, multi_set**, ...
 - ...
 - **string**
- So do other standard resources
 - **thread, lock_guard**, ...
 - **istream, fstream**, ...
 - **unique_ptr, shared_ptr**
 - ...



OOP




Class hierarchies



- Protection model
- No universal base class
 - an unnecessary implementation-oriented artifact
 - imposes avoidable space and time overheads.
 - encourages underspecified (overly general) interfaces
- Multiple inheritance
 - Separately consider interface and implementation
 - Abstract classes provide the most stable interfaces
- Minimal run-time type identification
 - **dynamic_cast<D*>(pb)**
 - **typeid(p)**

Inheritance

- Use it
 - When the domain concepts are hierarchical
 - When there is a need for run-time selection among hierarchically ordered alternatives



- Warning:
 - Inheritance has been seriously and systematically overused and misused
 - “When your only tool is a hammer everything looks like a nail”



GP



Generic Programming: Templates

- 1980: Use macros to express generic types and functions
- 1987 (and current) aims:
 - Extremely general/flexible
 - “must be able to do much more than I can imagine”
 - Zero-overhead
 - vector/Matrix/... to compete with C arrays
 - Well-specified interfaces
 - Implying overloading, good error messages, and maybe separate compilation
- “two out of three ain’t bad”
 - But it isn’t really good either
 - it has kept me concerned/working for 20+ years

Templates

- Compile-time duck typing
 - Leading to template metaprogramming
- A massive success in C++98, better in C++11, better still in C++14
 - STL containers
 - **template<typename T> class vector { /* ... */ };**
 - STL algorithms
 - **sort(v.begin(),v.end());**
 - And much more
- Better support for compile-time programming
 - C++11: **constexpr** (improved in C++14)

Algorithms

- Messy code is a major source of errors and inefficiencies
- We must use more explicit, well-designed, and tested algorithms
- The C++ standard-library algorithms are expressed in terms of half-open sequences [**first:last**)
 - For generality and efficiency

```
void f(vector<int>& v, list<string>& lst)
{
    sort(v.begin(),v.end());                // sort the vector using <

    auto p = find(lst.begin(),lst.end(),"Aarhus"); // find "Aarhus" in the list

    // ...
}
```

- We parameterize over element type and container type

Algorithms

- Simple, efficient, and general implementation
 - For any forward iterator
 - For any (matching) value type

```
template<typename Iter, typename Value>  
Iter find(Iter first, Iter last, Value val) // find first p in [first:last) so that *p==val  
{  
    while (first!=last && *first!=val)  
        ++first;  
    return first;  
}
```

Algorithms and Function Objects

- Parameterization with criteria, actions, and algorithms
 - Essential for flexibility and performance

```
void g(vector< string>& vs)
{
    auto p = find_if(vs.begin(), vs.end(), Less_than{"Griffin"});

    // ...
}
```


Algorithms and Function Objects

- The implementation is still trivial

```
template<typename Iter, typename Predicate>
Iter find_if(Iter first, Iter last, Predicate pred) // find first p in [first:last) so that pred(*p)
{
    while (first!=last && !pred(*first))
        ++first;
    return first;
}
```

Function Objects and Lambdas

- General function object
 - Can carry state
 - Easily inlined (i.e., close to optimally efficient)

```
struct Less_than {  
    String s;  
    Less_than(const string& ss) :s{ss} {} // store the value to compare against  
    bool operator()(const string& v) const { return v<s; } // the comparison  
};
```

Lambda notation

- We can let the compiler write the function object for us

```
auto p = std::find_if(vs.begin(),vs.end(),  
                    [](const string& v) { return v<"Griffin"; } );
```

Container algorithms

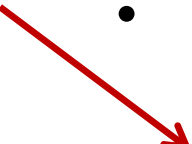
- The C++ standard-library algorithms are expressed in terms of half-open sequences [first:last)
 - For generality and efficiency
 - If you find that verbose, define container algorithms

```
namespace Extended_STL {  
    // ...  
    template<typename C, typename Predicate>  
    Iterator<C> find_if(C& c, Predicate pred)  
    {  
        return std::find_if(c.begin(),c.end(),pred);  
    }  
    // ...  
}
```

```
auto p = find_if(v, [](int x) { return x%2; } );    // assuming v is a vector<int>
```

Duck Typing is Insufficient



- There are no proper interfaces
 - Leaves error detection far too late
 - Compile- and link-time in C++
 - Encourages a focus on implementation details
 - Entangles users with implementation
 - Leads to over-general interfaces and data structures
 - As programmers rely on exposed implementation “details”
 - Does not integrate well with other parts of the language
 - Teaching and maintenance problems
 - We must think of generic code in ways similar to other code
 - Relying on well-specified interfaces (like OO, etc.)
- 

Generic Programming is just Programming

- *Traditional code*

```
double sqrt(double d);    // C++84: accept any d that is a double  
double d = 7;  
double d2 = sqrt(d);      // fine: d is a double  
double d3 = sqrt(&d);     // error: &d is not a double
```

- *Generic code*

```
void sort(Container& c);   // C++14: accept any c that is a Container  
vector<string> vs { "Hello", "new", "World" };  
sort(vs);                 // fine: vs is a Container  
sort(&vs);                 // error: &vs is not a Container
```

C++14: Constraints aka “Concepts lite”

- How do we specify requirements on template arguments?
 - state intent
 - Explicitly states requirements on argument types
 - provide point-of-use checking
 - No checking of template definitions
 - use constexpr functions
- Voted as C++14 Technical Report
- Design by B. Stroustrup, G. Dos Reis, and A. Sutton
- Implemented by Andrew Sutton in GCC
- There are no C++0x concept complexities
 - No concept maps
 - No new syntax for defining concepts
 - No new scope and lookup issues

What is a Concept?

- Concepts are fundamental
 - They represent fundamental concepts of an application area
 - Concepts are come in “clusters” describing an application area
- A concept has semantics (meaning)
 - Not just syntax
 - “**Subtractable**” is not a concept
- We have always had concepts
 - C++: Integral, arithmetic
 - STL: forward iterator, predicate
 - Informally: Container, Sequence
 - Algebra: Group, Ring, ...



What is a Concept?

- Don't expect to find a new fundamental concept every year
- A concept is ***not*** the minimal requirements for an implementation
 - An implementation does not define the requirements
 - Requirements should be stable
- Concepts support interoperability
 - There are relatively few concepts
 - We can remember a concept



C++14 Concepts (Constraints)

- A concept is a predicate on one or more arguments

- E.g. **Sequence<T>()** *// is T a Sequence?*

- Template declaration

```
template <typename S, typename T>  
    requires Sequence<S>()  
           && Equality_comparable<Value_type<S>, T>()  
Iterator_of<S> find(S& seq, const T& value);
```

- Template use

```
void use(vector<string>& vs)  
{  
    auto p = find(vs, "Jabberwocky");  
    // ...  
}
```

C++14 Concepts: Error handling

- Error handling is simple (and fast)

```
template<Sortable Cont>
    void sort(Cont& container);
```

```
vector<double> vec {1.2, 4.5, 0.5, -1.2};
list<int> lst {1, 3, 5, 4, 6, 8, 2};
```

```
sort(vec);      // OK: a vector is Sortable
sort(lst);      // Error at (this) point of use: Sortable requires random access
```

- **Actual** error message
error: 'list<int>' does not satisfy the constraint 'Sortable'

C++14 Concepts: “Shorthand Notation”

- Shorthand notation

```
template <Sequence S, Equality_comparable<Value_type<S>> T>  
    Iterator_of<C> find(S& seq, const T& value);
```

- We can handle essentially all of the Palo Alto TR
 - (STL algorithms) and more
 - Except for the axiom parts
 - We see no problems checking template definitions in isolation
 - But proposing that would be premature (needs work, experience)
 - We don't need explicit **requires** much (the shorthand is usually fine)

C++14 Concepts: Overloading

- Overloading is easy

```
template <Sequence S, Equality_comparable<Value_type<S>> T>  
    Iterator_of<S> find(S& seq, const T& value);
```

```
template<Associative_container C>  
    Iterator_type<C> find(C& assoc, const Key_type<C>& key);
```

```
vector<int> v { /* ... */ };
```

```
multiset<int> s { /* ... */ };
```

```
auto vi = find(v, 42);
```

// calls 1st overload:

// a vector is a Sequence

```
auto si = find(s, 12-12-12);
```

// calls 2nd overload:

// a multiset is an Associative_container

C++14 Concepts: Overloading

- Overloading based on predicates
 - specialization based on subset
 - Far easier than writing lots of tests

```
template<Input_iterator Iter>
```

```
    void advance(Iter& p, Difference_type<Iter> n) { while (n--) ++p; }
```

```
template<Bidirectional_iterator Iter>
```

```
    void advance(Iter& i, Difference_type<Iter> n)
```

```
    { if (n > 0) while (n--) ++p; if (n < 0) while (n++) --ip }
```

```
template<Random_access_iterator Iter>
```

```
    void advance(Iter& p, Difference_type<Iter> n) { p += n; }
```

- We don't say

```
    Input_iterator < Bidirectional_iterator < Random_access_iterator
```

we compute it



C++14 Concepts: Definition

- How do you write constraints?
 - Any **bool** expression
 - Including type traits and constexpr function
 - a **requires(expr)** expression
 - **requires()** is a compile time intrinsic function
 - **true** if **expr** is a valid expression
- To recognize a concept syntactically, we can declare it **concept**
 - Rather than just **constexpr**

C++14 Concepts: “Terse Notation”

- We can use a concept name as the name of a type than satisfy the concept

void sort(Container& c); *// terse notation*

– means

template<Container __Cont> *// shorthand notation*
void sort(__Cont& c);

– means

template<typename __Cont> *// explicit use of predicate*
requires Container<__Cont>()
void sort(__Cont)& c;

– Accepts any type that is a Container

vector<string> vs;
sort(vs);

C++14 Concepts: “Terse Notation”

- We have reached the conventional notation
 - with the conventional meaning

- *Traditional code*

```
double sqrt(double d);    // C++84: accept any d that is a double
double d = 7;
double d2 = sqrt(d);      // fine: d is a double
double d3 = sqrt(&d);     // error: &d is not a double
```

- *Generic code*

```
void sort(Container& c);   // C++14: accept any c that is a Container
vector<string> vs { "Hello", "new", "World" };
sort(vs);                 // fine: vs is a Container
sort(&vs);                // error: &vs is not a Container
```


C++14 Concepts: “Terse Notation”

- Consider `std::merge`
- Explicit use of predicates:

```
template<typename For,  
        typename For2,  
        typename Out>  
requires Forward_iterator<For>()  
        && Forward_iterator<For2>()  
        && Output_iterator<Out>()  
        && Assignable<Value_type<For>,Value_type<Out>>()  
        && Assignable<Value_type<For2>,Value_type<Out>>()  
        && Comparable<Value_type<For>,Value_type<For2>>()  
void merge(For p, For q, For2 p2, For2 q2, Out p);
```

- Headache inducing, and `accumulate()` is worse

C++14 Concepts: “Terse Notation”


- Better, use the shorthand notation

```
template<Forward_iterator For,  
        Forward_iterator For2,  
        Output_iterator Out>
```

```
    requires Mergeable<For,For2,Out>()
```

```
    void merge(For p, For q, For2 p2, For2 q2, Out p);
```

- Quite readable



C++14 Concepts: “Terse Notation”

- Better still, use the “terse notation”:

```
Mergeable{For,For2,Out}  // Mergeable is a concept requiring three types  
void merge(For p, For q, For2 p2, For2 q2, Out p);
```

- The

concept-name { identifier-list }

notation introduces constrained names

- Make simple things simple!

C++14 Concepts: “Terse Notation”

- Now we just need to define **Mergeable**:

```
template<typename For, typename For2, typename Out>
concept bool Mergeable()
{
    return Forward_iterator<For>()
        && Forward_iterator<For2>()
        && Output_iterator<Out>()
        && Assignable<Value_type<For>,Value_type<Out>>>()
        && Assignable<Value_type<For2>,Value_type<Out>>>()
        && Comparable<Value_type<For>,Value_type<For2>>>();
}
```

- It's just a predicate

Challenges



C++ Challenges

- Obviously, C++ is not perfect
 - How can we make programmers prefer modern styles over low-level code
 - which is far more error-prone and harder to maintain, yet no more efficient?
 - How can we make C++ a better language given the Draconian constraints of C and C++ compatibility?
 - How can we improve and complete the techniques and models (incompletely and imperfectly) embodied in C++?
- Solutions that eliminate major C++ strengths are not acceptable
 - Compatibility
 - link, source code
 - Performance
 - uncompromising
 - Portability
 - Range of application areas
 - Preferably increasing the range

Long-term C++ Challenges

- Close more type loopholes
 - in particular, find a way to prevent misuses of **delete** without spoiling RAI
- Simplify concurrent programming
 - in particular, provide some higher-level concurrency models as libraries
- Simplify generic programming
 - in particular, introduce simple and effective concepts
- Simplify programming using class hierarchies
 - in particular, eliminate use of the visitor pattern
- Better support for combinations of object-oriented and generic programming
- Make exceptions usable for hard-real-time projects
 - that will most likely be a tool rather than a language change
- Find a good way of using multiple address spaces
 - as needed for distributed computing
 - would probably involve defining a more general module mechanism that would also address dynamic linking, and more.
- Provide many more domain-specific libraries
- Develop a more precise and formal specification of C++


“Paradigms”

- Much of the distinction between object-oriented programming, generic programming, and “conventional programming” is an illusion
 - based on a focus on language features
 - incomplete support for a synthesis of techniques
 - The distinction does harm
 - by limiting programmers, forcing workarounds

```
void draw_all(Container& c)    // is this OOP, GP, or conventional?  
    requires Same_type<Value_type<Container>,Shape*>  
{  
    for_each(c, [](Shape* p) { p->draw(); } );  
}
```


Questions?

C++: A light-weight abstraction programming language



Key strengths:

- software infrastructure
- resource-constrained applications

Practice type-rich programming